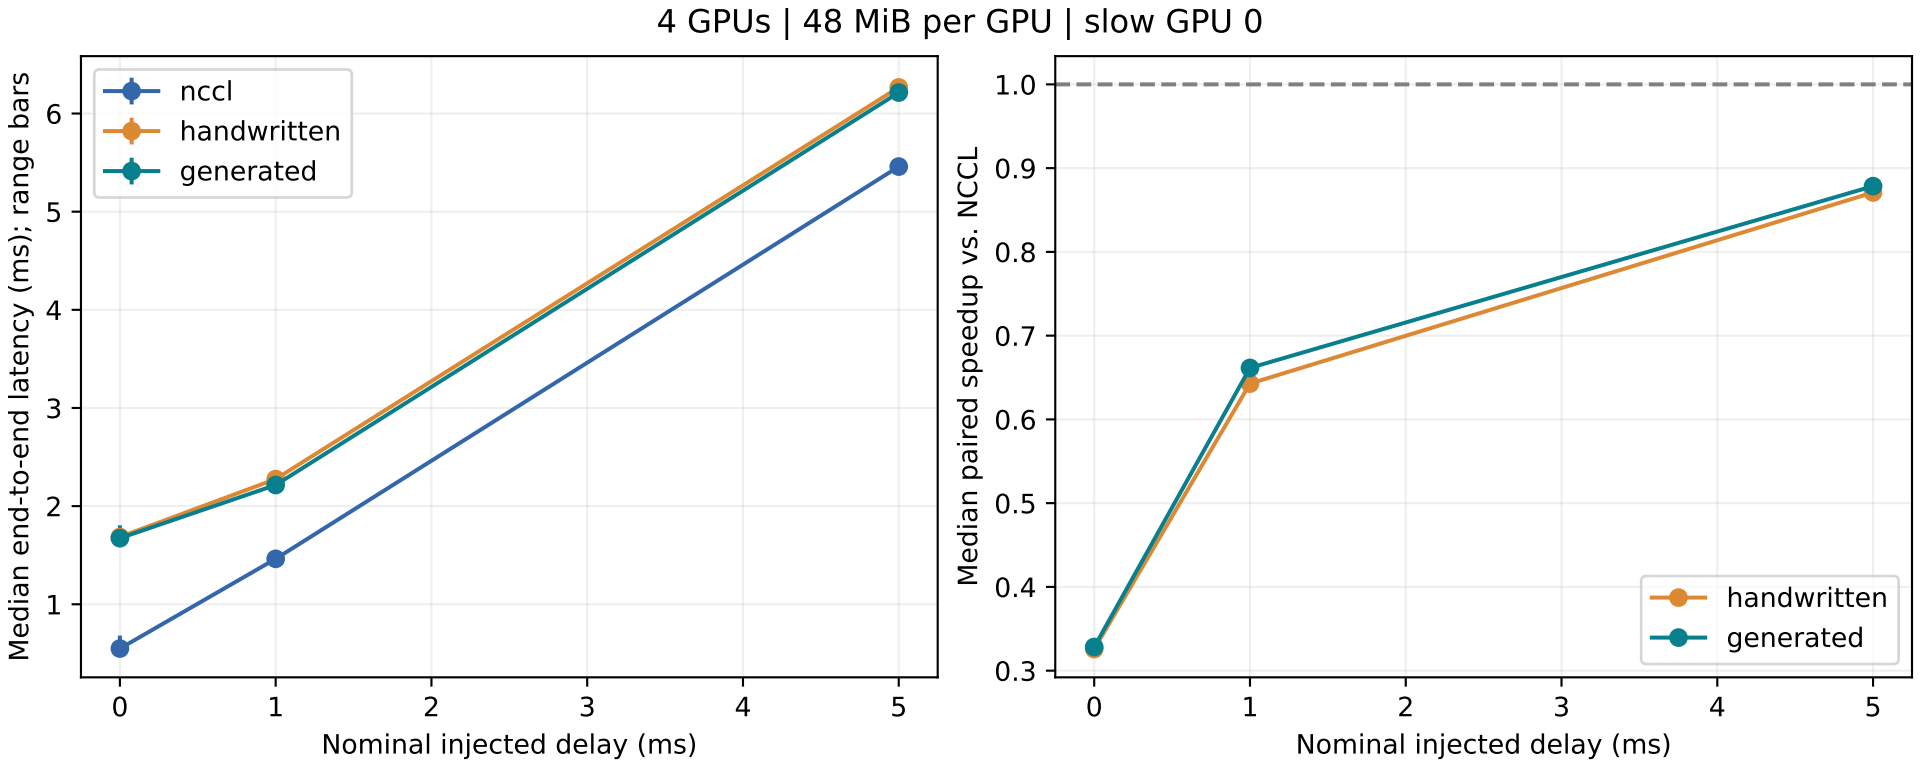
4 GPUs | 48 MiB per GPU | slow GPU 0
nccl
handwritten
generated
Median end-to-end latency (ms); range bars
Nominal injected delay (ms)
Median paired speedup vs. NCCL
handwritten
generated
Nominal injected delay (ms)